

Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](http://vision.stanford.edu/teaching/cs231n/assignments.html) (<http://vision.stanford.edu/teaching/cs231n/assignments.html>) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
In [1]: import random
import numpy as np
from cs231n.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

```
In [2]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
        """
        Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
        it for the linear classifier. These are the same steps as we used
        for the SVM, but condensed to a single function.
        """
        # Load the raw CIFAR-10 data
        cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'

        # Cleaning up variables to prevent loading data multiple times (which may cause memory issue)
```

```

try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

```

```

X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

```

```

# subsample the data

```

```

mask = list(range(num_training, num_training + num_validation))
X_val = X_train[mask]
y_val = y_train[mask]
mask = list(range(num_training))
X_train = X_train[mask]
y_train = y_train[mask]
mask = list(range(num_test))
X_test = X_test[mask]
y_test = y_test[mask]
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]

```

```

# Preprocessing: reshape the image data into rows

```

```

X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

```

```

# Normalize the data: subtract the mean image

```

```

mean_image = np.mean(X_train, axis = 0)
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

```

```

# add bias dimension and transform into columns

```

```

X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

```

```

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_de

```

v

```

# Invoke the above function to get our data.

```

```

X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIF
AR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)

```

```
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)
```

```
Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)
```

Softmax Classifier

Your code for this section will all be written inside `cs231n/classifiers/softmax.py`.

```
In [8]: # First implement the naive softmax loss function with nested loops.
        # Open the file cs231n/classifiers/softmax.py and implement the
        # softmax_loss_naive function.

        from cs231n.classifiers.softmax import softmax_loss_naive
        import time

        # Generate a random softmax weight matrix and use it to compute the loss.
        W = np.random.randn(3073, 10) * 0.0001
        loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

        # As a rough sanity check, our loss should be something close to -log(0.1).
        print('loss: %f' % loss)
        print('sanity check: %f' % (-np.log(0.1)))
```

```
loss: 2.380704
sanity check: 2.302585
```

Inline Question 1

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.**

Your Answer : Our initialization number of classes is low, which means that the probability predictions will all huddle around the same distribution at initialization, close to 0.1.

```
In [9]: # Complete the implementation of softmax_loss_naive and implement a (n
aive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging
tool.
# The numeric gradient should be close to the analytic gradient.
from cs231n.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

```
numerical: -0.714321 analytic: -0.714321, relative error: 2.023323e-
08
numerical: -3.132088 analytic: -3.132088, relative error: 2.194385e-
08
numerical: -0.416735 analytic: -0.416735, relative error: 3.867904e-
08
numerical: 1.327177 analytic: 1.327177, relative error: 2.587317e-08
numerical: 1.399740 analytic: 1.399740, relative error: 4.753285e-08
numerical: 1.812898 analytic: 1.812898, relative error: 2.535014e-08
numerical: -0.888163 analytic: -0.888163, relative error: 2.589727e-
08
numerical: -1.257603 analytic: -1.257603, relative error: 3.617679e-
08
numerical: -1.854077 analytic: -1.854077, relative error: 1.288557e-
08
numerical: -1.224152 analytic: -1.224153, relative error: 2.582007e-
08
numerical: 1.321453 analytic: 1.321453, relative error: 4.461755e-09
numerical: -2.744751 analytic: -2.744751, relative error: 6.710593e-
09
numerical: 0.364409 analytic: 0.364409, relative error: 1.493146e-08
numerical: -3.907450 analytic: -3.907450, relative error: 4.122210e-
09
numerical: 2.346587 analytic: 2.346587, relative error: 9.298321e-09
numerical: -0.118372 analytic: -0.118372, relative error: 6.468934e-
08
numerical: -0.007598 analytic: -0.007598, relative error: 4.914527e-
06
numerical: 0.591932 analytic: 0.591932, relative error: 5.035477e-08
numerical: 0.348856 analytic: 0.348856, relative error: 1.091617e-07
numerical: -1.033043 analytic: -1.033043, relative error: 2.639177e-
08
```

```
In [16]: # Now that we have a naive implementation of the softmax loss function
         # and its gradient,
         # implement a vectorized version in softmax_loss_vectorized.
         # The two versions should compute the same results, but the vectorized
         # version should be
         # much faster.
         tic = time.time()
         loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
         toc = time.time()
         print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

         from cs231n.classifiers.softmax import softmax_loss_vectorized
         tic = time.time()
         loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y
         _dev, 0.000005)
         toc = time.time()
         print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc -
         tic))

         # As we did for the SVM, we use the Frobenius norm to compare the two
         # versions
         # of the gradient.
         grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fr
         o')
         print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
         print('Gradient difference: %f' % grad_difference)

naive loss: 2.380704e+00 computed in 0.106275s
vectorized loss: 0.000000e+00 computed in 0.002561s
Loss difference: 2.380704
Gradient difference: 0.000000
```

```
In [25]: # Use the validation set to tune hyperparameters (regularization stren
         # gth and
         # learning rate). You should experiment with different ranges for the
         # learning
         # rates and regularization strengths; if you are careful you should be
         # able to
         # get a classification accuracy of over 0.35 on the validation set.

         from cs231n.classifiers import Softmax
         results = {}
         best_val = -1
         best_softmax = None

         #####
         #####
         # TODO:
         #
         # Use the validation set to set the learning rate and regularization s
         trength. #
         # This should be identical to the validation that you did for the SVM;
```

```

save      #
# the best trained softmax classifier in best_softmax.
#
#####

# Provided as a reference. You may or may not want to change these hyperparameters
learning_rates = [1e-7, 5e-7]
regularization_strengths = [2.5e3, 5e3]

# *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

for lr in learning_rates:
    for reg in regularization_strengths:
        softmax = Softmax() #specifying function
        softmax.train(X_train, y_train, learning_rate = lr, reg = reg,
num_iters = 1500)

        #train
        y_train_pred = softmax.predict(X_train)
        train_accuracy = np.mean(y_train == y_train_pred)

        #validation
        y_val_pred = softmax.predict(X_val)
        val_accuracy = np.mean(y_val == y_val_pred)

        if val_accuracy > best_val:
            best_val = val_accuracy
            best_softmax = softmax

        results[(lr, reg)] = train_accuracy, val_accuracy

        ##updating reg strength down by one factor (4-3) gave a 3 point
jump in cross valid accuracy.

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f'
% best_val)

```

```
lr 1.000000e-07 reg 2.500000e+03 train accuracy: 0.274163 val accuracy: 0.276000
lr 1.000000e-07 reg 5.000000e+03 train accuracy: 0.298265 val accuracy: 0.294000
lr 5.000000e-07 reg 2.500000e+03 train accuracy: 0.389429 val accuracy: 0.384000
lr 5.000000e-07 reg 5.000000e+03 train accuracy: 0.385673 val accuracy: 0.393000
best validation accuracy achieved during cross-validation: 0.393000
```

```
In [26]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))
```

```
softmax on raw pixels final test set accuracy: 0.373000
```

Inline Question 2 - True or False

Suppose the overall training loss is defined as the sum of the per-datapoint loss over all training examples. It is possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your Answer : yes, it is possible - True.

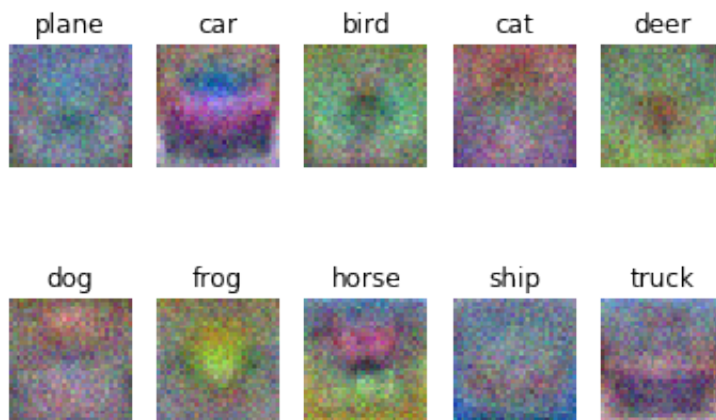
Your Explanation : with softmax, because we are in log scale, the overall value added to our classifier will always be greater than 0.

```
In [27]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias
w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])
```



In []: